

LEDE PROGRAM: DATA AND DATABASES DAY 2

Here's my summary yesterday's SQL fun in class.

First we listed the tables using the \d command.

```
mondial=# \d
```

Schema	Name	Type	Owner
public	airport	table	postgres
public	borders	table	postgres
public	city	table	postgres
public	citylocalname	table	postgres
public	cityothername	table	postgres
public	citypops	table	postgres
public	continent	table	postgres
public	country	table	postgres
public	countrylocalname	table	postgres
public	countryothername	table	postgres
public	countrypops	table	postgres
public	desert	table	postgres
public	economy	table	postgres
public	encompasses	table	postgres
public	ethnicgroup	table	postgres
public	geo_desert	table	postgres
public	geo_estuary	table	postgres
public	geo_island	table	postgres
public	geo_lake	table	postgres
public	geo_mountain	table	postgres
public	geo_river	table	postgres
public	geo_sea	table	postgres
public	geo_source	table	postgres
public	island	table	postgres
public	islandin	table	postgres
public	ismember	table	postgres
public	lake	table	postgres
public	lakeonisland	table	postgres
public	language	table	postgres
public	located	table	postgres
public	locatedon	table	postgres
public	mergeswith	table	postgres
public	mountain	table	postgres
public	mountainonisland	table	postgres
public	organization	table	postgres
public	politics	table	postgres
public	population	table	postgres
public	province	table	postgres
public	provincelocalname	table	postgres
public	provinceothername	table	postgres
public	provpops	table	postgres
public	religion	table	postgres
public	river	table	postgres
public	riveronisland	table	postgres
public	riverthrough	table	postgres
public	sea	table	postgres

(46 rows)

Note that on a Mac control + f and control + b lets you page down or page up respectively-instead of pressing return or the down arrow to get one line at a time.

Also note that you can exit the list by pressing q. This will get you back to the prompt without having to view every single result.

Next we checked the columns of the table "mountainonisland" using \d tablename

```
mondial=# \d mountainonisland
```

Column	Type	Collation	Nullable	Default
mountain	character varying(50)		not null	
island	character varying(50)		not null	

Indexes:

"mountainisland" PRIMARY KEY, btree (mountain, island)

Then we selected all the columns and rows in that table.

```
mondial=# SELECT * FROM mountainisland;
```

mountain	island
Gunnbjörn Fjeld	Greenland
Newtontoppen	Svalbard
Hvannadalshnukur	Iceland
Snaefell	Iceland
Hekla	Iceland
Katla	Iceland
Higravstinden	Austvágá, y
Ben Nevis	Great Britain
Snowdon	Great Britain
Sgàrr Alasdair	Skye
Carrauntoohil	Ireland
Slieve Donard	Ireland
Puig Major	Mallorca
Monte Cinto	Corse
Punta La Marmora	Sardegna
Monte Capanne	Elba
Etna	Sicilia
Pizzo Carbonara	Sicilia
Aenos	Kefallinia
Elati	Lefkas
Dirfi	Euboea
Fengari	Samothraki
Kerkis	Samos
Pilineo	Chios
Pramnos	Ikaria
Attavyros	Rhodos
Psiloritis	Crete
Olympos	Cyprus
Pico de Teide	Teneriffa

We decided to order the output by island to see if anything interesting appeared.

```
mondial=# SELECT * FROM mountainisland ORDER BY island;
```

mountain	island
Salahutu	Ambon
Ntringui	Anjouan
Higravstinden	Austvágá, y
Mt. Odin	Baffin Island
Agung	Bali
La Soufriere	Basse-Terre
Pico Basile	Bioko
Besar	Borneo
Bukit Raya	Borneo
Bukit Batubrok	Borneo
Murud	Borneo
Siho	Borneo
Mantam	Borneo
Kinabalu	Borneo
Mt. Balbi	Bougainville
Kapalatmada	Buru
Osmeña Peak	Cebu
Binaiya	Ceram
Pilineo	Chios
Monte Cinto	Corse
Psiloritis	Crete
Pico Turquino	Cuba
Olympos	Cyprus
Treuter Mt.	Devon Island
Morne Diablotins	Dominica
Monte Capanne	Elba
Barbeau Peak	Ellesmere Island
Tabwemasana	Espiritu Santo
Dirfi	Euboea

From the look of Borneo, there are some islands that have multiple mountains on them. So we rushed into aggregate queries and decided

to use COUNT and GROUP BY to count how many mountains each island had. **Note** that we counted the column 'island' to see how many rows each island appeared in.

```
mondial=# SELECT island, COUNT(island) FROM mountainonisland GROUP BY island;
```

island	count
Basse-Terre	1
Cyprus	1
Mallorca	1
Samosir	1
Unimak	1
Guadalcanal	1
Jamaica	1
Crete	1
Luzon	6
Lanai	1
Samos	1
Svalbard	1
Sao Jorge	1
Bougainville	1
Kolombangara	1
Grande Terre	1
Halmahera	1
Krenizyn	1
Chios	1
Isla da Ometepe	1
Sumatra	11
Cuba	1
Fogo	1
Hispaniola	2
Upolu	1
Krakatau	1
Panay	1
Gomera	1
Teresa Island	1

Next we tried to add the 'mountain' column to this query. But we got an error:

```
mondial=#
ERROR: column "mountainonisland.mountain" must appear in the GROUP BY clause or be used in an aggregate function
LINE 1: SELECT island, COUNT(island), mountain FROM mountainonisland...
```

The error tells us that you can't just put a normal column into an aggregate query.

So we went ahead and decided to get an ordered list of the number of mountains on each island, using ORDER BY

```
mondial=# SELECT island, COUNT(island) FROM mountainonisland GROUP BY island ORDER BY COUNT(island);
```

island	count
Faial	1
Cyprus	1
Mallorca	1
Samosir	1
Unimak	1
Guadalcanal	1
Jamaica	1
Crete	1
Kauai	1
Leyte	1
Oahu	1
Sardegna	1
Novaya Zemlya Yuzhny Island	1
Negros	1
Goodenough Island	1
Anjouan	1
Rhodos	1
Baffin Island	1

Because ORDER BY defaults to ascending values ASC we saw the beginning of a list that starts with all the mountains that just have one island. So we used DESC to change the order--and get the highest values first.

```
mondial=# SELECT island, COUNT(island) FROM mountainonisland GROUP BY island ORDER BY COUNT(island) DESC;
```

island	count
Sumatra	11
New Guinea	9
Java	7
Borneo	7
Sulawesi	7
Luzon	6
Honshu	6
Iceland	4
Mindanao	4
Madagaskar	3
Te Waka-a-Maui (South Island)	3
Hispaniola	2
Ireland	2
Flores	2
Te Ika-a-Maui (North Island)	2
Great Britain	2
Hawaii	2
Sicilia	2
Jeju	2
Reunion	2
Timor	2
Upolu	1
Krakatau	1
Panay	1
Gomera	1
Teresa Island	1
Terceira	1

Sumatra wins! At least according to the database. Next I introduced AS which allows you to rename columns (give them aliases). Below is the exact same query as above except that I renamed the column.

```
mondial=# SELECT island, COUNT(island) AS mostmn FROM mountainonisland GROUP BY island ORDER BY mostmn DESC;
```

island	mostmn
Sumatra	11
New Guinea	9
Java	7
Borneo	7
Sulawesi	7
Luzon	6
Honshu	6
Iceland	4
Mindanao	4
Madagaskar	3
Te Waka-a-Maui (South Island)	3
Hispaniola	2
Ireland	2
Flores	2
Te Ika-a-Maui (North Island)	2
Great Britain	2
Hawaii	2
Sicilia	2
Jeju	2
Reunion	2
Timor	2
Upolu	1
Krakatau	1
Panay	1
Gomera	1
Teresa Island	1
Terceira	1

Next, we jumped out of aggregate queries, and into the city table

```
mondial=# \d city
```

Table "public.city"				
Column	Type	Collation	Nullable	Default
name	character varying(50)		not null	
country	character varying(4)		not null	
province	character varying(50)		not null	
population	numeric			
latitude	numeric			

```

longitude | numeric
elevation | numeric
Indexes:
    "citykey" PRIMARY KEY, btree (name, country, province)
Check constraints:
    "citylat" CHECK (latitude >= '-90'::integer::numeric AND latitude <= 90::numeric)
    "citylon" CHECK (longitude >= '-180'::integer::numeric AND longitude <= 180::numeric)
    "citypop" CHECK (population >= 0::numeric)

```

We selected all of the columns to see what's in there.

```

mondial=# SELECT * FROM city;

```

name	country	province	population	latitude
Tirana	AL	Albania	418495	41.33
Shkodër	AL	Albania	77075	42.07
Durrës	AL	Albania	113249	41.32
Vlorë	AL	Albania	79513	40.47
Elbasan	AL	Albania	78703	41.1
Korçë	AL	Albania	51152	40.6
Komotini	GR	Anatolikos Makedonias kai Thrakis		41.1
Kavala	GR	Anatolikos Makedonias kai Thrakis	58790	40.93
Athina	GR	Attikis	664046	37.97
Peiraias	GR	Attikis	163688	37.95
Peristeri	GR	Attikis	139981	38.02
Acharnes	GR	Attikis	106943	38.08

A lot of columns, not so easy to read. So we specified our columns:

```

mondial=# SELECT name, country, population, elevation FROM city;

```

name	country	population	elevation
Tirana	AL	418495	110
Shkodër	AL	77075	13
Durrës	AL	113249	0
Vlorë	AL	79513	25
Elbasan	AL	78703	150
Korçë	AL	51152	850
Komotini	GR		45
Kavala	GR	58790	0
Athina	GR	664046	70
Peiraias	GR	163688	0
Peristeri	GR	139981	50
Acharnes	GR	106943	186
Patra	GR	213984	0
Kozani	GR		710
Kerkyra	GR		0
Ioannina	GR	112486	480
Thessaloniki	GR	325182	0
Iraklio	GR	173993	150
Chania	GR	108642	0
Ermoupoli	GR		0
Rhodes	GR	115490	26
Tripoli	GR		655
Lamia	GR	75315	50
Chalkida	GR	102223	0
Larissa	GR	162591	67
Volos	GR	144449	0
Mytilini	GR		8

Much easier to read! So we decided to look for cities at the lowest elevation using ORDER BY

```

mondial=# SELECT name, country, population, elevation FROM city ORDER BY elevation;

```

name	country	population	elevation
Baku	AZ	2150800	-28
Astrakhan	R	527345	-28
Atyrau	KAZ	196494	-20
Aktau	KAZ	181526	-8
David	PA	144858	-6
Almere	NL	196244	-3
Lelystad	NL	76252	-3

Georgetown	GUY	118363	-2
New Orleans	USA	343829	-2
Babol	IR	250217	-2
Bergen	N	267950	0
Malabo	GQ	92900	0
Kerkyra	GR		0
Stockholm	S	881235	0
Thessaloniki	GR	325182	0
Chania	GR	108642	0
Ermoupoli	GR		0
Delhi	IND	11034555	0
Chalkida	GR	102223	0
Callao	PE	438326	0
Volos	GR	144449	0
Saint-Paul	REUN	103916	0
Brest	F	140547	0
Ajaccio	F	66245	0
Le Havre	F	174156	0
Praia	CV	131719	0
Marseille	F	850636	0

Then we decided to look for the cities that the highest elevation by adding DESC to reverse the order.

```
mondial=# SELECT name, country, population, elevation FROM city ORDER BY elevation DESC;
```

name	country	population	elevation
Shashi	CN	281352	
Victoria	SY	24970	
Jamestown	HELX		
Saint-Denis	REUN	145347	
Tekirdag	TR	150112	
Korla	CN	159344	
Karab�k	TR	110537	
Aksu	CN	164092	
Yalova	TR	102874	
Duzce	TR	135557	
Osmaniye	TR	209255	
Yining	CN	177193	
Shihezi	CN	299676	
Wulumuqi	CN	3029372	
Shizuishan	CN	257862	
Manzhouli	CN	120023	
Linhe	CN	133183	
Ulanhot	CN	159538	
Jining	CN	163552	
Sor�	DK		
Wuzhou	CN	210452	
T��rshavn	FARX	13130	
Mariehamn	SF	10851	
H��meenlinna	SF	67803	
Lahti	SF	103396	
Tampere	SF	220678	
Kuopio	SF	106475	
Lappeenranta	SF	72617	

But something weird happened! We got all of these rows with no values for the elevation. This is because there were null values (missing values), and you have to tell the database what to do with them or they will show up somewhere. So we added NULLS LAST to the ORDER BY command, to make sure that they get pushed to the end of the list.

```
mondial=# SELECT name, country, population, elevation FROM city ORDER BY elevation DESC NULLS LAST;
```

name	country	population	elevation
Cerro de Pasco	PE		4330
Lhasa	CN	106885	4200
El Alto	BOL	848840	4150
Potos�	BOL	189652	4067
Puno	PE	120229	3830
Juliaca	PE	216716	3825
Oruro	BOL	264683	3735
Huancavelica	PE		3676
La Paz	BOL	764617	3640
Cusco	PE	348935	3399
Huancayo	PE	323054	3259

Huaraz	PE		100931		3052
Quito	EC		1619146		2850
Tunja	CO		184864		2820
Sucre	BOL		259388		2810
Ayacucho	PE		151019		2761
Cajamarca	PE		162326		2750
Sacaba	BOL		169494		2719
Toluca	MEX		489333		2660
Bogotá	CO		7776845		2640
Soacha	CO		500097		2566
Cuenca	EC		331888		2560
Cochabamba	BOL		630587		2558
Pasto	CO		434486		2527
Dessie	ETH		147592		2470
Zacatecas	MEX		129011		2440
Pachuca	MEX		256584		2432

Then we got into aggregates again, and compared the two different aggregate commands COUNT() and SUM()

```
mondial=# SELECT COUNT(population) FROM city;
count
-----
  3051
(1 row)

mondial=# SELECT SUM(population) FROM city;
sum
-----
1762114278
(1 row)
```

There are a few interesting things here. COUNT() is adding up unique values in the column, not numbers. Most cities have different populations, but there are some that have the same number and a bunch that have missing entries. If you do a count of all the cities (SELECT COUNT(name) FROM city;) you will find that there are 3390 unique cities, as opposed to 3051 unique population numbers. SUM() is actually adding up all the numbers.

Also note that you don't need GROUP BY if you just want aggregate all of the values in a single column.

Then we added up the city populations for each country. Here we need GROUP BY.

```
mondial=# SELECT country, SUM(population) FROM city GROUP BY country;
country | sum
-----+-----
NEP      | 1699953
RA       | 15351698
CH       | 1381946
L        | 99852
AMSA     |
WD       | 14725
LB       | 1010970
BF       | 1965190
EW       | 502624
MV       | 133019
BEN      | 1378292
CAYM     |
SY       | 24970
IRQ      | 12950153
GROX     |
NOK      | 9087275
MNG      | 760077
MW       | 1335704
AFG      | 3371100
NZ       | 1222122
BI       | 497166
WAFU     |
WG       | 37000
J        | 48415803
I        | 14398216
SK       | 1106091
ARU      |
```

Now we want to know what the actual countries are, because these country codes are quite cryptic. So we need to look for the country table to get the full name.

```
mondial=# \d country
```

```
Table "public.country"
  Column      |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
name         | character varying(50) |          | not null |
code        | character varying(4)  |          | not null |
capital     | character varying(50) |          |          |
province    | character varying(50) |          |          |
area        | numeric               |          |          |
population  | numeric               |          |          |
```

```
Indexes:
```

```
    "countrykey" PRIMARY KEY, btree (code)
```

```
    "country_name_key" UNIQUE CONSTRAINT, btree (name)
```

```
Check constraints:
```

```
    "countryarea" CHECK (area >= 0::numeric)
```

```
    "countrypop" CHECK (population >= 0::numeric)
```

```
mondial=# SELECT * FROM country;
```

name	code	capital	province	area	pop
Albania	AL	Tirana	Albania	28750	
Greece	GR	Athina	Attikis	131940	:
Macedonia	MK	Skopje	Macedonia	25333	
Serbia	SRB	Beograd	Serbia	77474	
Montenegro	MNE	Podgorica	Montenegro	14026	
Kosovo	KOS	Prishtine	Kosovo	10887	
Andorra	AND	Andorra la Vella	Andorra	450	
France	F	Paris	Île-de-France	547030	
Spain	E	Madrid	Madrid	504750	,
Austria	A	Wien	Wien	83850	
Czech Republic	CZ	Praha	Praha	78703	:
Germany	D	Berlin	Berlin	356910	{
Hungary	H	Budapest	Budapest	93030	

After a slight tangent investigating what the province column meant, we went ahead and did a JOIN between the two tables.

```
mondial=# SELECT SUM(city.population), country.name
```

```
mondial=# FROM city JOIN country ON country.code = city.country GROUP BY country.name;
```

sum	name
1145367	Costa Rica
703963	Cambodia
775000	Eritrea
46045206	Turkey
951418	Chad
207482	Cyprus
36735	Samoa
394368	Slovenia
	Cayman Islands
50182	Kiribati
14446981	Vietnam
637411	Kuwait
662426	Jamaica
22219	Antigua and Barbuda
5002435	Cameroon
10367222	Saudi Arabia
131719	Cape Verde
622712	Macedonia
1047048	Bosnia and Herzegovina
	Guinea-Bissau
125000	Curacao
5454065	Netherlands
	Svalbard
357238	Gambia
22867466	Zaire
49410900	Nigeria
1100000	Lebanon

Some things to note about JOIN: first you need to use DOT notation (city.population) to specify the columns you are selecting. The first part is name of the table (city.) The second part is the column from that table (.population).

JOIN is part of the FROM command. You put the first table after FROM then the second table after JOIN and then the criteria for joining (country.code = city.country) comes after ON. To put this in order that makes sense:

FROM firsttable JOIN secondtable ON condition (country.code = city.country)

Next we ordered the resulting joined table. Note that all the nulls came up first. But then the countries with large populations showed up after them.

```
SELECT SUM(city.population), country.name
FROM city JOIN country ON country.code = city.country
GROUP BY country.name
ORDER BY SUM(city.population) DESC;
```

sum	name
	Turks and Caicos Islands
	French Polynesia
	Nauru
	Jersey
	Cook Islands
	Macao
	Norfolk Island
	Wallis and Futuna
	British Virgin Islands
	Guernsey
	Somalia
	Northern Mariana Islands
	Cayman Islands
	Anguilla
	Isle of Man
	Greenland
	Guam
	American Samoa
	Falkland Islands
	Bermuda
	Cocos Islands
	Mayotte
	Gibraltar
	Guinea-Bissau
	Saint Helena
	Guadeloupe
	Svalbard
	Aruba
	Christmas Island
	Mauritius
	Niue
326058186	China
129752758	India
93687185	Brazil
81882315	United States
72000673	Russia
50920843	Pakistan
49410900	Nigeria
48415803	Japan
47137170	Mexico
46627466	Indonesia
46045206	Turkey
34845642	Iran
33570031	South Korea
26579496	Colombia
25333235	Germany

Again, we pushed down the null values.

```
FROM city JOIN country ON country.code = city.country GROUP BY country.name ORDER BY SUM(city.population) DESC NULLS FIRST;
```

sum	name
326058186	China
129752758	India
93687185	Brazil
81882315	United States
72000673	Russia
50920843	Pakistan
49410900	Nigeria
48415803	Japan
47137170	Mexico
46627466	Indonesia
46045206	Turkey
34845642	Iran
33570031	South Korea
26579496	Colombia
25333235	Germany
25252422	United Kingdom

```

23470701 | South Africa
22867466 | Zaire
22364857 | Egypt
19594294 | Philippines
18553641 | Spain
17198533 | Ukraine
16915048 | Australia
16481172 | Taiwan
15670355 | Canada
15351698 | Argentina
14607665 | Bangladesh

```

Then we moved into 'mountain'

```
mondial=# \d mountain
```

```

      Table "public.mountain"
  Column      |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 name         | character varying(50) |           | not null |
 mountains    | character varying(50) |           |          |
 elevation     | numeric          |           |          |
 type         | character varying(10) |           |          |
 coordinates  | geocoord         |           |          |

```

Indexes:

```
"mountainkey" PRIMARY KEY, btree (name)
```

Check constraints:

```
"mountaincoord" CHECK ((coordinates).latitude >= '-90'::integer::numeric AND (coordinates).latitude <= 90::nume
```

```
mondial=# SELECT * FROM mountain;
```

name	mountains	elevation	type	coordinates
Gunnbj�rn Fjeld		3694		(68.92,-29.9)
Newtontoppen		1713		(79.0,17.3)
Hvannadalshnukur		2110	volcanic	(64.1,-16.7)
Snaefell		1833	volcanic	(64.5,-15.2)
Hekla		1491	volcano	(64,-19.7)
Katla		1450	volcano	(63.6,-19.0)
Haltiatunturi	Scandinavian Mountains	1365		(69.3,21.27)
Kebnekaise	Scandinavian Mountains	2099		(67.9,18.5)
Sarektjokko	Scandinavian Mountains	2089		(67.43,17.71)
Higravstinden		1146		(68.36,14.79)
Galdhoeppig	Scandinavian Mountains	2469		(61.64,8.31)
Glittertind	Scandinavian Mountains	2465		(61.65,8.55)
Snoehetta	Scandinavian Mountains	2286		(62.32,9.27)
Ben Nevis	Grampians	1344		(56.8,-5.0)
Snowdon		1015		(53.08,-4.08)
Sg�rr Alasdair		993		(57.21,-6.22)
Carrauntoohil	Armorican Highlands	1041		(52.00,-9.74)
Slieve Donard		849		(54.18,-5.92)
Feldberg	Black Forest	1493		(47.87,8.00)
Brocken	Harz	1141		(51.8,10.6)
Grosser Arber	Bayrischer Wald	1456		(49.1,13)
Zugspitze	Alps	2963		(47.42,10.99)
Barre des Ecrins	Alps	4101		(44.92,6.36)
Mont Ventoux	Alps	1912		(44.18,5.28)
Gran Paradiso	Alps	4061		(45.5,7.25)
Mont Blanc	Alps	4808		(45.83,6.86)
Grand Combin	Alps	4314		(45.93,7.30)

We ordered it alphabetically, to see if there were ranges with multiple mountains. And we narrowed down our columns. Note that we ordered first alphabetically, and then by elevation.

```
mondial=# SELECT name, mountains, elevation FROM mountain ORDER BY mountains, elevation;
```

name	mountains	elevation
Mt. Marcy	Adirondacks	1629
Tahat	Ahaggar	3003
Fansipan	Ailao Shan	3143
Jebel Shams	Al Hajar Mountains	3009
Ayrybaba	Alai	3138
Tschimtarga	Alai	5489
Pik Tandykul	Alai	5544
Mt. Ratz	Alaska Boundary Range	3090
Mt. McKinley (Denali)	Alaska Range	6193
Mt. Redoubt	Aleutian Range	3108
Mont Ventoux	Alps	1912

Hochgolling	Alps	2862
Triglav	Alps	2864
Zugspitze	Alps	2963
Marmolata	Alps	3343
Grossglockner	Alps	3797
Piz Bernina	Alps	4048
Gran Paradiso	Alps	4061
Barre des Ecrins	Alps	4101
Finsteraarhorn	Alps	4274
Grand Combin	Alps	4314
Matterhorn	Alps	4478
Monte Rosa	Alps	4634
Mont Blanc	Alps	4808
Khâiten Peak	Altai	4374
Bjelucha	Altai	4506
Altun Shan Peak	Altyn-Tagh	5830

Next we used WHERE to only get mountains with elevations above 4000.

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation > 4000 ORDER BY mountains, elevation;
```

name	mountains	elevation

Tschimtaga	Alai	5489
Pik Tandykul	Alai	5544
Mt. McKinley (Denali)	Alaska Range	6193
Piz Bernina	Alps	4048
Gran Paradiso	Alps	4061
Barre des Ecrins	Alps	4101
Finsteraarhorn	Alps	4274
Grand Combin	Alps	4314
Matterhorn	Alps	4478
Monte Rosa	Alps	4634
Mont Blanc	Alps	4808
Khâiten Peak	Altai	4374
Bjelucha	Altai	4506
Altun Shan Peak	Altyn-Tagh	5830
Sulamutag Feng	Altyn-Tagh	6245
Monte San Valentin	Andes	4058
Maipo	Andes	5264
Licancabur	Andes	5920
Nevado El Plomo	Andes	6070
Marmolejo	Andes	6108
Tupungato	Andes	6550
Llullaillaco	Andes	6739
Monte Pissis	Andes	6795
Ojos del Salado	Andes	6893
Aconcagua	Andes	6962
Sâphan DaŸÄ±	Armenian Highlands	4058
Ararat	Armenian Highlands	5165

Then we stopped ordering alphabetically, and just by elevation.

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation > 4000 ORDER BY elevation;
```

name	mountains	elevation

Boundary Peak	Sierra Nevada California	4006
Otgon Tenger	Khangai Range	4008
Wheeler Peak	Rocky Mountains	4011
Mt. Waddington	Pacific Ranges	4019
Mt. Victoria	Owen Stanley Range	4038
Piz Bernina	Alps	4048
Monte San Valentin	Andes	4058
Sâphan DaŸÄ±	Armenian Highlands	4058
Gran Paradiso	Alps	4061
Aragaz	Lesser Caucasus	4090
Fako	Cameroon Line	4095
Kinabalu	Crocker Range	4096
Barre des Ecrins	Alps	4101
Kings Peak	Rocky Mountains	4123
Pacha Mama		4138
Mt. Boising	Finisterre Range	4150
Tubkhal	Atlas	4167
Mauna Loa	Hawaii	4170
Gannett Peak	Rocky Mountains	4207
Mauna Kea	Hawaii	4214

Tajumulco	Sierra Madre de Chiapas	4220
Finsteraarhorn	Alps	4274
Pikes Peak	Rocky Mountains	4302
Grand Combin	Alps	4314
Elgon	East African Rift	4321
Nevado de Colima	Sierra Volcanica Transversal	4330
Mt. Giluwe	Bismarck Range	4368

Then we changed the order to get the highest elevation first. Note that we don't need to deal with null values because the WHERE command is making sure we only get values above 4000.

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation > 4000 ORDER BY elevation DESC;
```

name	mountains	elevation
-----+-----+-----		
Mt. Everest	Himalaya	8848
K2	Karakorum	8611
Kangchendzonga	Himalaya	8586
Lhotse	Himalaya	8516
Makalu	Himalaya	8485
Cho Oyu	Himalaya	8188
Dhaulagiri	Himalaya	8167
Manaslu	Himalaya	8163
Nanga Parbat	Himalaya	8125
Annapurna	Himalaya	8091
Gasherbrum I	Karakorum	8080
Broad Peak	Karakorum	8051
Gasherbrum II	Karakorum	8034
Shishapangma	Himalaya	8027
Nanda Devi	Himalaya	7816
Batura Sar	Karakorum	7795
Namcha Barwa	Himalaya	7782
Saltoro Kangri	Karakorum	7742
Tirich Mir	Hindukush	7708
Gurla Mandhata	Himalaya	7694
Saser Kangri	Karakorum	7672
Kongur	Pamir	7649
Gangkhar Puensum	Himalaya	7570
Gongga Shan	Hengduan Shan	7556
Muztagh Ata	Pamir	7509
Pik Ismoil Somoni	Pamir	7495
Noshaq	Hindukush	7492
Pik Pobeda	Tian Shan	7439

Here I introduced a different condition in the WHERE command BETWEEN

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation 4000 AND 6000 ORDER BY elevation DESC;
```

name	mountains	elevation
-----+-----+-----		
Alto Toroni	Cordillera Volcanica	5982
Mt. Logan	Elias Range	5959
Alpamayo	Cordillera Blanca	5947
Licancabur	Andes	5920
Cotopaxi	Cordillera Occidental	5897
Kilimanjaro	East African Rift	5895
Hkakabo Razi	Hengduan Shan	5881
Ollagñe	Cordillera Volcanica	5870
Altun Shan Peak	Altyn-Tagh	5830
Kangze'gyai	Qiliang Shan	5808
Cayambe	Cordillera Occidental	5796
Pico Cristobal Colon	Sierra Nevada de Santa Marta	5775
Nevado del Huila	Cordillera Occidental	5750
Zapaleri	Cordillera de L�pez	5653
Elbrus	Kaukasus	5642
Citlaltep�tl (Pico de Orizaba)	Sierra Volcanica Transversal	5636

Here we used the LIMIT command to only show a certain number of rows. This shows the first 10:

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation BETWEEN 4000 AND 6000 ORDER BY elevation
```

name	mountains	elevation
-----+-----+-----		
Alto Toroni	Cordillera Volcanica	5982
Mt. Logan	Elias Range	5959
Alpamayo	Cordillera Blanca	5947

Licancabur	Andes		5920
Cotopaxi	Cordillera Occidental		5897
Kilimanjaro	East African Rift		5895
Hkakabo Razi	Hengduan Shan		5881
Ollagñe	Cordillera Volcanica		5870
Altun Shan Peak	Altyn-Tagh		5830
Kangze'gyai	Qiliang Shan		5808

(10 rows)

This also includes OFFSET to show the next 10 rows:

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE elevation BETWEEN 4000 AND 6000 ORDER BY elevation
```

name		mountains		elevation
-----+-----+-----				
Cayambe		Cordillera Occidental		5796
Pico Cristobal Colon		Sierra Nevada de Santa Marta		5775
Nevado del Huila		Cordillera Occidental		5750
Zapaleri		Cordillera de L�pez		5653
Elbrus		Kaukasus		5642
Citlaltep�tl (Pico de Orizaba)		Sierra Volcanica Transversal		5636
Damavand		Elburs		5610
Shanzidou		Hengduan Shan		5596
Qilian Shan Peak		Qiliang Shan		5547
Pik Tandykul		Alai		5544

(10 rows)

Here we used the IS NULL in the WHERE command to look for mountains with missing ranges.

```
mondial=# SELECT name, mountains, elevation FROM mountain WHERE mountains IS NULL ORDER BY elevation DESC LIMIT 10
```

name		mountains		elevation
-----+-----+-----				
Semeru				3676
Rantemario				3478
Baru				3475
Koryaksky				3456
Slamet				3428
Raung				3332
Etna				3323
Shiveluch				3307
Hasan Dagi				3268
Kita-Dake				3193

(10 rows)

Here we used AND in the WHERE command to look for mountains with missing ranges AND that had elevations higher than 4000.

```
SELECT name, mountains, elevation FROM mountain WHERE mountains IS NULL AND elevation > 4000 ORDER BY elevation DE
```

name		mountains		elevation
-----+-----+-----				
Kljutschewskaja Sopka				4750
Hazaran				4500
Pacha Mama				4138

Here is another aggregate query counting mountains in ranges.

```
mondial=# SELECT mountains, COUNT(mountains) from mountain GROUP BY mountains ORDER BY COUNT(mountains) DESC;
```

mountains		count
-----+-----		
Andes		19
Himalaya		15
Alps		14
Rocky Mountains		13
Barisan Mountains		12
East African Rift		11
Hawaii		7
Karakorum		7
Pamir		7
Dinaric Alps		6
Kaukasus		6
Scandinavian Mountains		6
Pyrenees		5
Cameroon Line		5

Canary Islands		5
Apennin		5
Cordillera Occidental		5
Azores		5
Sierra Volcanica Transversal		5
Tian Shan		5
Atlas		5
Ural		5
Cordillera Volcanica		5
Annamite Range		4
Elias Range		4
Hengduan Shan		4
Cascade Range		4

Here we added a new column to the aggregate query to get total elevations

```
SELECT mountains, COUNT(mountains), SUM(elevation) from mountain GROUP BY mountains ORDER BY COUNT(mountains) DESC,
```

mountains		count		sum
-----+-----+-----				
Andes		19		90721
Himalaya		15		121193
Alps		14		52459
Rocky Mountains		13		48021
Barisan Mountains		12		31423
East African Rift		11		45958
Hawaii		7		16798
Karakorum		7		55985
Pamir		7		50238
Dinaric Alps		6		13466
Kaukasus		6		28890
Scandinavian Mountains		6		12773
Pyrenees		5		15213
Cameroon Line		5		14560
Canary Islands		5		11081
Apennin		5		12092
Cordillera Occidental		5		29119
Azores		5		6573
Sierra Volcanica Transversal		5		25348
Tian Shan		5		28807
Atlas		5		12583
Ural		5		7736
Cordillera Volcanica		5		31107
Annamite Range		4		10542
Elias Range		4		21116
Hengduan Shan		4		25773
Cascade Range		4		14046

Here we used AVG() to get average elevations

```
SELECT mountains, COUNT(mountains), AVG(elevation) from mountain GROUP BY mountains ORDER BY COUNT(mountains) DESC,
```

mountains		count		avg
-----+-----+-----				
Andes		19		4774.7894736842105263
Himalaya		15		8079.5333333333333333
Alps		14		3747.0714285714285714
Rocky Mountains		13		3693.9230769230769231
Barisan Mountains		12		2618.5833333333333333
East African Rift		11		4178.0000000000000000
Hawaii		7		2399.7142857142857143
Karakorum		7		7997.8571428571428571
Pamir		7		7176.8571428571428571
Dinaric Alps		6		2244.3333333333333333
Kaukasus		6		4815.0000000000000000
Scandinavian Mountains		6		2128.8333333333333333
Pyrenees		5		3042.6000000000000000
Cameroon Line		5		2912.0000000000000000
Canary Islands		5		2216.2000000000000000
Apennin		5		2418.4000000000000000
Cordillera Occidental		5		5823.8000000000000000
Azores		5		1314.6000000000000000
Sierra Volcanica Transversal		5		5069.6000000000000000
Tian Shan		5		5761.4000000000000000
Atlas		5		2516.6000000000000000
Ural		5		1547.2000000000000000
Cordillera Volcanica		5		6221.4000000000000000

Annamite Range		4		2635.5000000000000000
Elias Range		4		5279.0000000000000000
Hengduan Shan		4		6443.2500000000000000
Cascade Range		4		3511.5000000000000000

We got a little sidetracked with geo_mountain, see below.

But first, to stick with this query, we went on to filter out averages below 3001 elevation (that is, selecting averages > 3000). We used HAVING which is like WHERE but for aggregates. HAVING must come after GROUP BY because it is a test on the aggregates.

```
mondial=# SELECT mountains, COUNT(mountains), AVG(elevation) from mountain GROUP BY mountains
mondial=# HAVING AVG(elevation) > 3000
mondial=# ORDER BY COUNT(mountains) DESC;
```

mountains		count		avg
Andes		19		4774.7894736842105263
Himalaya		15		8079.5333333333333333
Alps		14		3747.0714285714285714
Rocky Mountains		13		3693.9230769230769231
East African Rift		11		4178.0000000000000000
Pamir		7		7176.8571428571428571
Karakorum		7		7997.8571428571428571
Kaukasus		6		4815.0000000000000000
Tian Shan		5		5761.4000000000000000
Cordillera Occidental		5		5823.8000000000000000
Pyrenees		5		3042.6000000000000000
Cordillera Volcanica		5		6221.4000000000000000
Sierra Volcanica Transversal		5		5069.6000000000000000
Cascade Range		4		3511.5000000000000000
Hengduan Shan		4		6443.2500000000000000
Elias Range		4		5279.0000000000000000
Alai		3		4723.6666666666666667
Asir Mountains		3		3054.3333333333333333
Cordillera Blanca		3		6450.0000000000000000
Maoke Mountains		3		4798.0000000000000000
Armenian Highlands		3		4057.0000000000000000
Kunlun		3		7000.0000000000000000
Sierra Madre de Chiapas		3		3574.0000000000000000
Altai		2		4440.0000000000000000
Cordillera de Lpez		2		5830.5000000000000000
Lesser Kaukasus		2		3907.0000000000000000
Sredinny Range		2		3111.5000000000000000

You can have WHERE and HAVING in the same query. WHERE always comes first, and filters out the original table, before aggregation takes place. Also note that we used a different conditional LIKE() which looks for patterns in language.

```
mondial=# SELECT mountains, COUNT(mountains), AVG(elevation) from mountain
mondial=# WHERE mountains LIKE('A%')
mondial=# GROUP BY mountains
mondial=# HAVING AVG(elevation) > 3000
mondial=# ORDER BY COUNT(mountains) DESC;
```

mountains		count		avg
Andes		19		4774.7894736842105263
Alps		14		3747.0714285714285714
Asir Mountains		3		3054.3333333333333333
Alai		3		4723.6666666666666667
Armenian Highlands		3		4057.0000000000000000
Altai		2		4440.0000000000000000
Altyn-Tagh		2		6037.5000000000000000
Ahaggar		1		3003.0000000000000000
Ailao Shan		1		3143.0000000000000000
Arakan Mountains		1		3070.0000000000000000
Aleutian Range		1		3108.0000000000000000
Al Hajar Mountains		1		3009.0000000000000000
Alaska Boundary Range		1		3090.0000000000000000
Alaska Range		1		6193.0000000000000000

(14 rows)

Finally, here's the search we did for mountains in Greece! This joins the geo_mountain and mountains tables, searches for the country code, and orders by elevation.

```
mondial=# SELECT geo_mountain.country, mountain.name, mountain.elevation, mountain.type
mondial=# FROM geo_mountain JOIN mountain ON mountain.name = geo_mountain.mountain
```

```
mondial=# WHERE geo_mountain.country = 'GR'
mondial=# ORDER BY mountain.elevation DESC;
 country |      name      | elevation | type 
-----+-----+-----+-----
 GR      | Olymp          |      2917 | 
 GR      | Olymp          |      2917 | 
 GR      | Smolikas       |      2637 | 
 GR      | Profitis Ilias |      2497 | 
 GR      | Psiloritis     |      2456 | 
 GR      | Kyllini        |      2376 | 
 GR      | Athos          |      2033 | 
 GR      | Dirfi          |      1743 | 
 GR      | Aenos          |      1628 | 
 GR      | Fengari        |      1611 | 
 GR      | Kerkis         |      1433 | 
 GR      | Pilineo        |      1297 | 
 GR      | Attavyros      |      1215 | 
 GR      | Elati          |      1158 | 
 GR      | Pramnos        |      1037 | 
(15 rows)
```