

LEDE PROGRAM: DATA AND DATABASES DAY 2

Here's my summary yesterday's SQL fun in class.

First we listed the tables using the \d command.

```
mondial=# \d
```

List of relations			
Schema	Name	Type	Owner
public	airport	table	postgres
public	borders	table	postgres
public	city	table	postgres
public	citylocalname	table	postgres
public	cityothername	table	postgres
public	citypops	table	postgres
public	continent	table	postgres
public	country	table	postgres
public	countrylocalname	table	postgres
public	countryothername	table	postgres
public	countrypops	table	postgres
public	desert	table	postgres
public	economy	table	postgres
public	encompasses	table	postgres
public	ethnicgroup	table	postgres
public	geo_desert	table	postgres
public	geo_estuary	table	postgres
public	geo_island	table	postgres
public	geo_lake	table	postgres
public	geo_mountain	table	postgres
public	geo_river	table	postgres
public	geo_sea	table	postgres
public	geo_source	table	postgres
public	island	table	postgres
public	islandin	table	postgres
public	ismember	table	postgres
public	lake	table	postgres
public	lakeonisland	table	postgres
public	language	table	postgres
public	located	table	postgres
public	locatedon	table	postgres
public	mergeswith	table	postgres
public	mountain	table	postgres
public	mountainonisland	table	postgres
public	organization	table	postgres
public	politics	table	postgres
public	population	table	postgres
public	province	table	postgres
public	provincelocalname	table	postgres
public	provinceothername	table	postgres
public	provpops	table	postgres
public	religion	table	postgres
public	river	table	postgres
public	riveronisland	table	postgres
public	riverthrough	table	postgres
public	sea	table	postgres

(46 rows)

Note that on a Mac control + f and control + b lets you page down or page up respectively-instead of pressing return or the down arrow to get one line at a time. If you have a Windows machine and you have figured out how to, do this please post on slack!

Also note that you can exit the list by pressing q. This will get you back to the prompt without having to view every single result.

Next I showed you how leaving out the semicolon ; at the end will give you line after line of unexecuted commands until you finally put in the semicolon. Note that when the command prompt has an = you are at a fresh line (the first line of your query):

```
mondial=#
```

When the prompt has a - you are extending the query over more than one line:

mondial-#

```
mondial2=# SELECT * FROM mountainonisland
mondial2-#
mondial2-#
mondial2-#
mondial2-#
mondial2-#
mondial2-#
mondial2-# \
Invalid command \. Try \? for help.
mondial2-# SELECT *
mondial2-# FROM mountainonisland
mondial2-# ;
ERROR:  syntax error at or near "SELECT"
LINE 2: SELECT *
        ^
mondial2=# SELECT *
mondial2-# FROM mountainonisland
mondial2-# ;
```

Next we checked the columns of the table "mountainonisland" using \d tablename

```
mondial=# \d mountainonisland
              Table "public.mountainonisland"
  Column      |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 mountain     | character varying(50) |           | not null |
 island       | character varying(50) |           | not null |
Indexes:
    "mountainislkey" PRIMARY KEY, btree (mountain, island)
```

Then we selected all the columns and rows in that table.

```
mondial=# SELECT * FROM mountainonisland;
 mountain      | island
-----+-----
 Gunnbjörn Fjeld | Greenland
 Newtontoppen   | Svalbard
 Hvannadalshnukur | Iceland
 Snaefell       | Iceland
 Hekla          | Iceland
 Katla          | Iceland
 Higravstinden  | Austvågsey
 Ben Nevis      | Great Britain
 Snowdon        | Great Britain
 Sgairr Alasdair | Skye
 Carrauntoohil  | Ireland
 Slieve Donard  | Ireland
 Puig Major     | Mallorca
 Monte Cinto    | Corse
 Punta La Marmora | Sardegna
 Monte Capanne  | Elba
 Etna           | Sicilia
 Pizzo Carbonara | Sicilia
 Aenos         | Kefallinia
 Elati          | Lefkas
 Dirfi          | Euboea
 Fengari        | Samothraki
 Kerkis         | Samos
 Pilineo        | Chios
 Pramnos        | Icaria
 Attavyros      | Rhodos
 Psiloritis     | Crete
 Olympos        | Cyprus
 Pico de Teide   | Teneriffa
 Pico de las Nieves | Gran Canaria
 Pico de Malpaso | Hierro
 Garajonay      | Gomera
 Roque de los Muchachos | La Palma
 Pico Ruivo     | Madeira
 Pico           | Pico
```

Pico da Vara		Sao Miguel
Pico da Esperanãsa		Sao Jorge
Cabeãso Gordo		Faial
Serra de Santa Bãrbara		Terceira
Fogo		Fogo
Pico da Antãnia		Santiago
Pico Basile		Bioko
Pico de Sao Tome		Sao Tome
Queen Marys Peak		Tristan Da Cunha
Jabal Hajhir		Sokotra
Tsaratanana		Madagaskar
Tsiafajavona		Madagaskar
Andringitra		Madagaskar
Piton des Neiges		Reunion
Piton de la Fournaise		Reunion
Karthala		Grand Comoro
Ntringui		Anjouan
Pik Sedova		Novaya Zemlya Severny Island
Gora Pervousmotrennaya		Novaya Zemlya Yuzhny Island
Pidurutalagala		Sri Lanka
Chikurachki		Paramuschir
Krenizyn		Krenizyn
Lopatin		Sachalin
Zhima		Olkhon
Asahi-Dake		Hokkaido
Iwate		Honshu
Hotaka-Dake		Honshu
Kita-Dake		Honshu
Haku-San		Honshu
Fuji-San		Honshu
Daisen		Honshu
Ishizuchi-San		Shikoku
Sobo-San		Kyushu

Next we used WHERE to find only the mountains on the island of Honshu. Note, just because we see six islands in this list of rows, that doesn't mean another mountain on Honshu isn't hiding somewhere else because the list is not ordered.

```
mondial2=# SELECT * FROM mountainonisland WHERE island = 'Honshu';
 mountain | island 
-----+-----
 Iwate     | Honshu
 Hotaka-Dake | Honshu
 Kita-Dake  | Honshu
 Haku-San   | Honshu
 Fuji-San   | Honshu
 Daisen     | Honshu
(6 rows)
```

Then we added ORDER BY to get our mountains in alphabetical order (that's the kind of order computers like to do with text).

```
mondial2=# SELECT * FROM mountainonisland WHERE island = 'Honshu' ORDER BY mountain;
 mountain | island 
-----+-----
 Daisen   | Honshu
 Fuji-San | Honshu
 Haku-San | Honshu
 Hotaka-Dake | Honshu
 Iwate     | Honshu
 Kita-Dake | Honshu
(6 rows)
```

Finally we added LIMIT to our query, to show only the first two lines of the resulting table. Not very useful here, but it does keep things very simple.

```
mondial2=# SELECT * FROM mountainonisland WHERE island = 'Honshu' ORDER BY mountain LIMIT 2;
 mountain | island 
-----+-----
 Daisen   | Honshu
 Fuji-San | Honshu
(2 rows)
```

Then we jumped into a much more complex and messy (especially visually) table: city.

```
mondial2=# SELECT * FROM city;
```

name	country	province	population	latitude
Tirana	AL	Albania	418495	41.33
Shkodër	AL	Albania	77075	42.07
Durrës	AL	Albania	113249	41.32
Vlorë	AL	Albania	79513	40.47
Elbasan	AL	Albania	78703	41.1
Korçë	AL	Albania	51152	40.6
Komotini	GR	Anatolikos Makedonias kai Thrakis		41.1
Kavala	GR	Anatolikos Makedonias kai Thrakis	58790	40.93
Athina	GR	Attikis	664046	37.97
Peiraias	GR	Attikis	163688	37.95
Peristeri	GR	Attikis	139981	38.02
Acharnes	GR	Attikis	106943	38.08
Patra	GR	Dytikos Elladas	213984	38.25
Kozani	GR	Dytikos Makedonias		:

Here we learned the limits of using `SELECT *` in the command line. If there are a lot of columns it quickly becomes unreadable. So we took a look at the columns using `\d`, so we could decide what columns we wanted to select.

```
mondial2=# \d city
```

Table "public.city"				
Column	Type	Collation	Nullable	Default
name	character varying(50)		not null	
country	character varying(4)		not null	
province	character varying(50)		not null	
population	numeric			
latitude	numeric			
longitude	numeric			
elevation	numeric			

Indexes:

```
"citykey" PRIMARY KEY, btree (name, country, province)
```

Check constraints:

```
"citylat" CHECK (latitude >= '-90'::integer::numeric AND latitude <= 90::numeric)
```

```
"citylon" CHECK (longitude >= '-180'::integer::numeric AND longitude <= 180::numeric)
```

```
"citypop" CHECK (population >= 0::numeric)
```

We selected our columns, and also filtered (WHERE) so that we would get populations above 1 million. Getting a much more readable table.

```
mondial2=# SELECT name, country, province, population FROM city WHERE population > 1000000;
```

name	country	province	population
Beograd	SRB	Serbia	1639121
Paris	F	Île-de-France	2249975
Barcelona	E	Catalunya	1611013
Madrid	E	Madrid	3198645
Wien	A	Wien	1761738
Praha	CZ	Praha	1289556
München	D	Bayern	1348335
Berlin	D	Berlin	3292365
Hamburg	D	Hamburg	1706696
Köln	D	Nordrhein-Westfalen	1005775
Budapest	H	Budapest	1729040
Milano	I	Lombardia	1242123
Roma	I	Lazio	2617175
Minsk	BY	Minsk City	1836808
Warszawa	PL	Mazowieckie	1711324
Kharkiv	UA	Kharkivska	1441362
Odesa	UA	Odeska	1008162
Kyïv	UA	Kyïv	2814258
Sankt Peterburg	R	Sankt-Peterburg	5028000
Moskva	R	Moscow	11979529
Nizhnii Novgorod	R	Nizhnii Novgorodskaya	1259921
Voronezh	R	Voronezhskaya	1003638
Kazan	R	Tatarstan	1176187
Volgograd	R	Volgogradskaya	1018790
Samara	R	Samarskaya	1171598
Rostov-na-Donu	R	Rostovskaya	1103733
Ufa	R	Bashkortostan	1077719
Perm	R	Permskij	1013887
Yekaterinburg	R	Sverdlovskaya	1396074

```
mondial2=# SELECT name, country, province, population FROM city WHERE population > 1000000
mondial2=# ORDER BY population;
```

So we used DESC to show the highest populations first.

name	country	province	population
Shanghai	CN	Shanghai	22315474
Karachi	PK	Sindh	14916456
Lagos	WAN	Lagos	13745000
Istanbul	TR	Åstanbul	13710512
Mumbai	IND	Maharashtra	12442373
Moskva	R	Moscow	11979529
Beijing	CN	Beijing	11716620
Kinshasa	ZRE	Kinshasa	11575000
SÃ£o Paulo	BR	SÃ£o Paulo	11152344
Lahore	PK	Punjab	11126285
Tianjin	CN	Tianjin	11090314
Guangzhou	CN	Guangdong	11071424
Delhi	IND	Delhi	11034555
Shenzhen	CN	Guangdong	10358381
Seoul	ROK	South Korea	9805506
Wuhan	CN	Hubei	9785388
Jakarta	RI	DKI Jakarta	9607787
Tehran	IR	Tehran	8693706
Tokyo	J	Tokyo	8591695
Ciudad de MÃ©xico	MEX	Distrito Federal	8555272
Al Qahirah	ET	Egypt	8471859
Bangalore	IND	Karnataka	8443675
New York	USA	New York	8405837
Bangkok	THA	Thailand	8305218
London	GB	London	8250205
Dongguan	CN	Guangdong	8220207
BogotÃ;	CO	Santa Fe de BogotÃ;	7776845
Lima	PE	Lima City	7605742
Chongqing	CN	Chongqing	7457600

name	country	province	population
Ocoatepeque	HCA	Ocoatepeque	

Jamestown	HELX	Saint Helena
Mamoutzou	MAYO	Mayotte
Port Louis	MS	Mauritius
Sorðlög	DK	Sjælland
Otsu	J	Shiga
Nara	J	Nara
Labuan	MAL	Labuan
Anau	TM	Akhal
Guelmim	MA	Guelmim Es Semara
Zouerate	RIM	Tiris Zemmour
Tidjikja	RIM	Tagant
Wakayama	J	Wakayama
Akjoujt	RIM	Inchiri
Aioun	RIM	Hodh El Gharbi
Saint Peter Port	GBG	Guernsey
Nema	RIM	Hodh Chargui
Gibraltar	GBZ	Gibraltar
Selibaby	RIM	Guidimagha
Aleg	RIM	Brakna

A big difference! This is because in the previous query WHERE was not just filtering out low values, but also null values (empty fields in the columns). In psql if you order by descending, it defaults to putting the nulls first. So conveniently there is an additional command NULLS LAST .

```
mondial2=# SELECT name, country, province, population FROM city ORDER BY population DESC NULLS LAST;
```

name	country	province	population
Shanghai	CN	Shanghai	22315474
Karachi	PK	Sindh	14916456
Lagos	WAN	Lagos	13745000
Istanbul	TR	İstanbul	13710512
Mumbai	IND	Maharashtra	12442373
Moskva	R	Moscow	11979529
Beijing	CN	Beijing	11716620
Kinshasa	ZRE	Kinshasa	11575000
São Paulo	BR	São Paulo	11152344
Lahore	PK	Punjab	11126285
Tianjin	CN	Tianjin	11090314
Guangzhou	CN	Guangdong	11071424
Delhi	IND	Delhi	11034555
Shenzhen	CN	Guangdong	10358381
Seoul	ROK	South Korea	9805506
Wuhan	CN	Hubei	9785388

Next we used LIMIT again to just show the top 10.

```
mondial2=# SELECT name, country, province, population FROM city ORDER BY population DESC NULLS LAST
mondial2=# LIMIT 10;
```

name	country	province	population
Shanghai	CN	Shanghai	22315474
Karachi	PK	Sindh	14916456
Lagos	WAN	Lagos	13745000
Istanbul	TR	İstanbul	13710512
Mumbai	IND	Maharashtra	12442373
Moskva	R	Moscow	11979529
Beijing	CN	Beijing	11716620
Kinshasa	ZRE	Kinshasa	11575000
São Paulo	BR	São Paulo	11152344
Lahore	PK	Punjab	11126285

(10 rows)

We added WHERE to just get cities in China.

```
mondial2=# SELECT name, country, province, population FROM city WHERE country = 'CN' ORDER BY population DESC NULLS LAST;
```

name	country	province	population
Shanghai	CN	Shanghai	22315474
Beijing	CN	Beijing	11716620
Tianjin	CN	Tianjin	11090314
Guangzhou	CN	Guangdong	11071424
Shenzhen	CN	Guangdong	10358381
Wuhan	CN	Hubei	9785388
Dongguan	CN	Guangdong	8220207
Chongqing	CN	Chongqing	7457600

```

Chengdu | CN | Sichuan | 7415590
Foshan | CN | Guangdong | 7194311
(10 rows)

```

Next we add AND to WHERE have two parameters to filter our search for cities in China. We change the order to the default (ASC) so our resulting table begins with the least populous city that has more than 1 million people.

```

mondial2=# SELECT name, country, province, population FROM city
mondial2=# WHERE country = 'CN' AND population > 1000000 ORDER BY population;

```

name	country	province	population
Datong	CN	Shanxi	1110000
Handan	CN	Hebei	1110000
Luoyang	CN	Henan	1190000
Huainan	CN	Anhui	1200000
Jilin	CN	Jilin	1270000
Fushun	CN	Liaoning	1350000
Qiqihar	CN	Heilongjiang	1380000
Anshan	CN	Liaoning	1390000
Weifang	CN	Shandong	2044028
Baotou	CN	Nei Mongol	2096851
Zaozhuang	CN	Shandong	2125481
Xiangyang	CN	Hubei	2199690
Yantai	CN	Shandong	2227733
Nantong	CN	Jiangsu	2274113
Linyi	CN	Shandong	2303648
Huizhou	CN	Guangdong	2344634
Nanchang	CN	Jiangxi	2357839
Lanzhou	CN	Gansu	2628426
Huaian	CN	Jiangsu	2635406
Shijiazhuang	CN	Hebei	2834942
Fuzhou	CN	Fujian	2921762
Wulumuqi	CN	Xinjiang	3029372

I am skipping the very confusing example where I use OR instead of AND because it gives us very complicated results. The thing to understand with all conditional statements in all programming languages: with AND, the test has to be true for both cases, with OR the test only needs to be true for one of the cases, so you will get more results.

Here I introduced a few more of the possible conditional clauses starting with BETWEEN

```

mondial2=# SELECT name, country, province, population FROM city
mondial2=# WHERE population BETWEEN 100000 AND 1000000 ORDER BY population;

```

name	country	province	population
Fort-de-France	MART	Martinique	100000
Andria	I	Puglia	100052
Barakaldo	E	Euskadi	100064
Noginsk	R	Moskovskaya	100072
Francistown	RB	Botswana	100079
Kaspijsk	R	Dagestan	100129
Worcester	GB	West Midlands	100153
Lincoln	GB	East Midlands	100160
Piacenza	I	Emilia-Romagna	100311
Haining	CN	Zhejiang	100478
Masjed Soleyman	IR	Khuzestan	100497
Ancona	I	Marche	100497
SzĀkesfehĀrvĀr	H	FejĀr	100570
Giresun	TR	Giresun	100712
Huaraz	PE	Ancash	100931

IN() searches a list of values for specific matches. The search below is the same thing as "WHERE population=100000 OR population = 1000000" it is just much more efficient, especially if you wanted to put in a much longer list of values.

```

mondial2=# SELECT name, country, province, population FROM city
mondial2=# WHERE population IN (100000, 1000000) ORDER BY population;

```

name	country	province	population
Fort-de-France	MART	Martinique	100000

(1 row)

Finally we used LIKE() which uses regular expressions to search for word patterns to find country codes that begin with the letter F.

BD		13
BDS		1
BEN		4
BERM		1
BF		2
BG		6
BHT		1
BI		2
BIH		7
BOL		13
BR		210
BRN		1
BRU		1
BS		1
EVIR		1

Noting that 'BR' (Brazil) had 210 cities I did the following query to show you how COUNT() works to aggregate.

```
mondial2=# SELECT country FROM city WHERE country = 'BR';
country
```

[illegible]BR
(210 rows)

If we scroll to the end of this list we would see that there are 210 rows. So technically in this count, we are really just getting the number of rows in which 'BR' appears in the 'country' column.

Next, we ordered our aggregate count() so we could see which countries appear the most.

```
mondial2=# SELECT country, COUNT(country) FROM city GROUP BY country ORDER BY count(country) DESC;
```

country	count
-----+-----	
CN	302
USA	251
BR	210
R	171
IND	99
TR	88
D	85
GB	84
MEX	83
J	72
E	65
WAN	58
RI	58
I	56
CO	52
IR	50
RO	42
PL	41
F	41
RA	39
UA	38
CDN	37
RP	36
PK	35
YV	34
ZRE	33
PE	31
S	30

This prompted a discussion of the order of processes in psql queries. ORDER BY sorts the final table resulting from the query, after the grouping in aggregation is done.

But!!!! And this is where things can get confusing, we added WHERE to our aggregate query. And things got very different. Why? Because WHERE is not part of the aggregation. It filters the original table before the grouping and aggregation happens. So when we take out cities that have populations below 1000000, those rows disappear from the table that will be grouped and aggregated.

```
mondial2=# SELECT country, COUNT(country) FROM city WHERE population > 1000000
```

```
mondial2=# GROUP BY country ORDER BY count(country) DESC;
```

country	count
-----+-----	
CN	60
IND	44
R	15
BR	14
J	12
RI	11
PK	10
TR	10
WAN	10
USA	9
MEX	9
ROK	9
IR	8
RSA	6
RC	5
ZRE	5
AUS	5
RP	4
CO	4
ET	4
D	4
UA	3
CDN	3
MYA	3
RA	3
IRQ	3
YV	3
SA	3 (6 rows)

Importantly, I then introduced HAVING which is a conditional parameter--it tests and filters values. It uses the same conditional operations (AND, OR, IN(), etc) as WHERE. But instead of WHERE, which tests and filters values from the original table, HAVING filters the resulting table after it has been aggregated.

```
mondial2=# SELECT country, COUNT(country) FROM city WHERE population > 1000000
mondial2=# GROUP BY country HAVING count(country) > 9 ORDER BY count(country) DESC;
 country | count
-----+-----
 CN      |    60
 IND     |    44
 R       |    15
 BR      |    14
 J       |    12
 RI      |    11
 TR      |    10
 PK      |    10
 WAN     |    10
(9 rows)
```

You cannot have HAVING without GROUP BY . HAVING tests/filters that group.

There was a little digression on how ORDER BY works. Note the two queries. It shifts the final table. You can pass multiple columns, and it will sort starting with the first column and then the next.

```
mondial2=# SELECT country, COUNT(country) FROM city WHERE population > 1000000
 country | count
-----+-----
 WAN     |    10
 TR      |    10
 RI      |    11
 R       |    15
 PK      |    10
 J       |    12
 IND     |    44
 CN      |    60
 BR      |    14
(9 rows)
```

GROUP BY

```
mondial2=# SELECT country, COUNT(country) FROM city WHERE population > 1000000
 country | count
-----+-----
 CN      |    60
 IND     |    44
 R       |    15
 BR      |    14
 J       |    12
 RI      |    11
 PK      |    10
 TR      |    10
 WAN     |    10
(9 rows)
```

GROUP BY

Here I did another aggregate query using SUM() on the population column -- because it is a numerical column, SUM() works. It is adding up all of the populations by country. (You could try using AVG() to get the average city populations.)

Note, in the first query we get a lot of empty values because the sums were null.

```
mondial2=# SELECT country, SUM(population) FROM city
mondial2=# GROUP BY country ORDER BY SUM(population) DESC;
 country | sum
-----+-----
 GROX    |
 NAU     |
 GNB     |
 FPOL    |
 NORF    |
 COOK    |
 XMAS    |
 CAYM    |
 GBM     |
 MACX    |
 MS      |
 GBZ     |
 BVIR    |
```

```

AMSA |
AXA  |
MAYO |
SVAX |
ARU  |
SP   |
GBG  |
GUAD |
GUAM |
HELX |
FALK |
GBJ  |
TUCA |
NIUE |
BERM |
mondial2=# SELECT country, SUM(population) FROM city
GROUP BY country ORDER BY SUM(population) DESC NULLS LAST;
country |      sum
-----+-----
CN      | 326058186
IND     | 129752758
BR      | 93687185
USA     | 81882315
R       | 72000673
PK      | 50920843
WAN     | 49410900
J       | 48590545
MEX     | 47137170
RI      | 46893102
TR      | 46045206
IR      | 34845642
ROK     | 33570031
CO      | 26579496
D       | 25333235
GB      | 25252422
RSA     | 23470701
ZRE     | 22867466
ET      | 22364857
RP      | 19594294
E       | 18553641
UA      | 17198533
AUS     | 16915048
RC      | 16481172
CDN     | 15670355
RA      | 15351698
BD      | 14607665
VN      | 14446981

```

Another small digression/clarification on SELECT led me to this query, which just shows the sums and does not show the related countries because be left out the country column.

```

mondial2=# SELECT SUM(population) FROM city
mondial2=# GROUP BY country ORDER BY SUM(population) DESC NULLS LAST;
sum
-----
326058186
129752758
93687185
81882315
72000673
50920843
49410900
48590545
47137170
46893102
46045206
34845642
33570031
26579496
25333235
25252422
23470701
22867466
22364857
19594294
18553641
17198533

```

16915048
16481172
15670355
15351698
14607665
14446981

Finally we are on our way to using JOIN so that we can see the actual names of the countries. But first we want to look at the 'country' table, because that is the only table that contains the full names of the countries.

```
mondial2=# \d country
```

Table "public.country"				
Column	Type	Collation	Nullable	Default
name	character varying(50)		not null	
code	character varying(4)		not null	
capital	character varying(50)			
province	character varying(50)			
area	numeric			
population	numeric			

Indexes:

```
"countrykey" PRIMARY KEY, btree (code)
"country_name_key" UNIQUE CONSTRAINT, btree (name)
```

Check constraints:

```
"countryarea" CHECK (area >= 0::numeric)
"countrypop" CHECK (population >= 0::numeric)
```

The key column, the main column we would use to join is country.code (in all the other tables that key is called 'country' (like city.country)). Once we are doing a JOIN it is important to use dot notation (table.column) to specify which table each column is coming from.

Here we take our previous aggregate query, but we use JOIN to get the full name of the country: country.name -- NOTE: that JOIN is inside the FROM parameter. Until now FROM has been very simple because we just want one table. But once you start joining you are making hybrid tables using multiple pre-existing tables. These tables are brought together using ON which matches values from key columns that are shared across the different tables.

It might be helpful to imagine this entire statement: 'FROM city JOIN country ON city.country = country.code' to be the same thing as a single new customized table that you have invented.

```
mondial2=# SELECT country.name, SUM(city.population) FROM city
mondial2=# JOIN country ON city.country = country.code
mondial2=# GROUP BY country.name ORDER BY SUM(city.population) DESC NULLS LAST;
```

name	sum
China	326058186
India	129752758
Brazil	93687185
United States	81882315
Russia	72000673
Pakistan	50920843
Nigeria	49410900
Japan	48590545
Mexico	47137170
Indonesia	46893102
Turkey	46045206
Iran	34845642
South Korea	33570031
Colombia	26579496
Germany	25333235
United Kingdom	25252422
South Africa	23470701
Zaire	22867466
Egypt	22364857
Philippines	19594294
Spain	18553641
Ukraine	17198533
Australia	16915048
Taiwan	16481172
Canada	15670355
Argentina	15351698
Bangladesh	14607665
Vietnam	14446981

m

Below I breakdown how JOIN works to create a brand-new table. Below shows a single table 'city', then that table with the new column that was brought over from 'country'. This is what happens before the aggregation.

```
mondial2=# SELECT country, population FROM city ORDER BY country;
```

country	population
A	146676
A	269211
A	12046
A	124386
A	59239
A	193511
A	1761738
A	45922
A	27831
A	59942
A	96531
A	13485
A	52100
AFG	335200
AFG	311800
AFG	288700
AFG	2435400
AG	22219
AL	51152
AL	418495
AL	77075
AL	113249
AL	79513
AL	78703
AMSA	
AND	22256
ANG	
ANG	

```
mondial2=# SELECT city.country, city.population, country.name
```

```
mondial2-# FROM city JOIN country ON city.country = country.code
```

```
mondial2-# ORDER BY city.country;
```

country	population	name
A	146676	Austria
A	269211	Austria
A	12046	Austria
A	124386	Austria
A	59239	Austria
A	193511	Austria
A	1761738	Austria
A	45922	Austria
A	27831	Austria
A	59942	Austria
A	96531	Austria
A	13485	Austria
A	52100	Austria
AFG	335200	Afghanistan
AFG	311800	Afghanistan
AFG	288700	Afghanistan
AFG	2435400	Afghanistan
AG	22219	Antigua and Barbuda
AL	51152	Albania
AL	418495	Albania
AL	77075	Albania
AL	113249	Albania
AL	79513	Albania
AL	78703	Albania
AMSA		American Samoa
AND	22256	Andorra
ANG		Angola
ANG		Angola

Here I just change the order of appearance of the columns the exact same table above.

```
mondial2=# SELECT city.country, country.name, city.population
```

```
mondial2-# FROM city JOIN country ON city.country = country.code
```

```
mondial2-# ORDER BY country;
```

country	name	population
A	Austria	146676

A	Austria	269211
A	Austria	12046
A	Austria	124386
A	Austria	59239
A	Austria	193511
A	Austria	1761738
A	Austria	45922
A	Austria	27831
A	Austria	59942
A	Austria	96531
A	Austria	13485
A	Austria	52100
AFG	Afghanistan	335200
AFG	Afghanistan	311800
AFG	Afghanistan	288700
AFG	Afghanistan	2435400
AG	Antigua and Barbuda	22219
AL	Albania	51152
AL	Albania	418495
AL	Albania	77075
AL	Albania	113249
AL	Albania	79513
AL	Albania	78703
AMSA	American Samoa	
AND	Andorra	22256
ANG	Angola	
ANG	Angola	

Same aggregate query as above, but this time I put the country names in alphabetical order--note this will give you a different order than alphabetizing by country codes.

```
mondial2=# SELECT country.name, SUM(city.population) FROM city
JOIN country ON city.country = country.code
GROUP BY country.name ORDER BY country.name;
```

name	sum
-----+-----	
Afghanistan	3371100
Albania	818187
Algeria	7235824
American Samoa	
Andorra	22256
Angola	7273439
Anguilla	
Antigua and Barbuda	22219
Argentina	15351698
Armenia	1066264
Aruba	
Australia	16915048
Austria	2862618
Azerbaijan	2763900
Bahamas	248948
Bahrain	143035
Bangladesh	14607665
Barbados	88529
Belarus	3878063
Belgium	2256798
Belize	67186
Benin	1378292
Bermuda	
Bhutan	42465
Bolivia	5139125
Bosnia and Herzegovina	1047048
Botswana	327412
Brazil	93687185

Two more examples of simple JOINS without aggregating. Noting that with SELECT you're just specifying which columns to actually show.

```
mondial2=# SELECT city.country, country.name, city.population
mondial2=# FROM city JOIN country ON city.country = country.code
mondial2=# ORDER BY country;
```

country	name	population
-----+-----		
A	Austria	146676
A	Austria	269211
A	Austria	12046

```

A      | Austria      | 124386
A      | Austria      | 59239
A      | Austria      | 193511
A      | Austria      | 1761738
A      | Austria      | 45922
A      | Austria      | 27831
A      | Austria      | 59942
A      | Austria      | 96531
A      | Austria      | 13485
A      | Austria      | 52100
AFG    | Afghanistan  | 335200
AFG    | Afghanistan  | 311800
AFG    | Afghanistan  | 288700
AFG    | Afghanistan  | 2435400
AG     | Antigua and Barbuda | 22219
AL     | Albania      | 51152
AL     | Albania      | 418495
AL     | Albania      | 77075
AL     | Albania      | 113249
AL     | Albania      | 79513
AL     | Albania      | 78703
AMSA   | American Samoa | 
AND    | Andorra      | 22256
ANG    | Angola       | 
ANG    | Angola       | 

```

```

mondial2=# SELECT city.country, country.code, country.name, city.population
mondial2=# FROM city JOIN country ON city.country = country.code
mondial2=# ORDER BY country;

```

country	code	name	population
A	A	Austria	146676
A	A	Austria	269211
A	A	Austria	12046
A	A	Austria	124386
A	A	Austria	59239
A	A	Austria	193511
A	A	Austria	1761738
A	A	Austria	45922
A	A	Austria	27831
A	A	Austria	59942
A	A	Austria	96531
A	A	Austria	13485
A	A	Austria	52100
AFG	AFG	Afghanistan	335200
AFG	AFG	Afghanistan	311800
AFG	AFG	Afghanistan	288700
AFG	AFG	Afghanistan	2435400
AG	AG	Antigua and Barbuda	22219
AL	AL	Albania	51152
AL	AL	Albania	418495
AL	AL	Albania	77075
AL	AL	Albania	113249
AL	AL	Albania	79513
AL	AL	Albania	78703
AMSA	AMSA	American Samoa	
AND	AND	Andorra	22256
ANG	ANG	Angola	
ANG	ANG	Angola	

Finally, I introduced AS which allows you to alias columns (and tables as well). You can give them more simple/clear names using AS, and you can use those aliases in specific parts of your query to make things more efficient. Note the column headers in the results are different.

```

mondial2=# SELECT city.country AS countryCode, country.name AS cName, city.population AS pop
mondial2=# FROM city JOIN country ON city.country = country.code
mondial2=# ORDER BY country;

```

countrycode	cname	pop
A	Austria	146676
A	Austria	269211
A	Austria	12046
A	Austria	124386
A	Austria	59239
A	Austria	193511
A	Austria	1761738
A	Austria	45922
A	Austria	27831

A	Austria	59942
A	Austria	96531
A	Austria	13485
A	Austria	52100
AFG	Afghanistan	335200
AFG	Afghanistan	311800
AFG	Afghanistan	288700
AFG	Afghanistan	2435400
AG	Antigua and Barbuda	22219
AL	Albania	51152
AL	Albania	418495
AL	Albania	77075
AL	Albania	113249
AL	Albania	79513
AL	Albania	78703
AMSA	American Samoa	
AND	Andorra	22256
ANG	Angola	
ANG	Angola	